

# Large Meaning Models (LMMs): Phaneron as a Language-Agnostic Structural Semantics Substrate

Anonymous Preprint

November 13, 2025

## Abstract

We introduce *Large Meaning Models* (LMMs): language systems whose core is an explicit, language-agnostic *meaning substrate*—the *Phaneron*—where distinctions and connections are primitive and all nuance (roles,  $n$ -ary relations, direction, tense/aspect, discourse links) is represented as reusable *patterns* and *graph rewrites*. We formalize an LMM as an abstract data type with (i) a versioned Phaneron meaning graph, (ii) bidirectional bridges between surface and meaning, and (iii) a hybrid objective that *compresses* discourse and *predicts* future utterances/world updates. We map core linguistic phenomena to unlabeled structural motifs, outline learning/inference via constraint-driven abduction and MDL+prediction, and propose evaluations for paraphrase invariance, cross-lingual alignment, explainable QA, and long-range discourse coherence. LMMs complement LLMs: LLMs act as string transducers; LMMs are the inspectable *meaning* core.

## 1 Introduction

Token-next architectures blur surface form with meaning, offer weak inspectability, wobble under paraphrase, and struggle with long-range discourse state. We propose to separate *meaning* from *surface* with a single-layer, unlabeled structural substrate where semantics live as reusable patterns discovered by a rewrite loop guided by compression and prediction.

## 2 Background: The Phaneron in brief

The Phaneron is an unlabeled, undirected pseudograph  $G = (V, E)$  with reusable *patterns* (finite subgraphs), double-pushout (DPO) graph rewrites, and a versioned event log  $\mathcal{E}$ . Meaning equals structural position (ideally automorphism orbits; operationally,  $r$ -hop/WL colors). A hybrid objective  $\mathcal{L} = \mathcal{L}_{\text{MDL}} + \lambda \mathcal{L}_{\text{pred}}$  scores states and candidate rewrites.

## 3 The LMM ADT

**Definition 1** (Large Meaning Model). *An LMM is a tuple  $\mathcal{M} = (\mathcal{P}, \mathcal{B}_{\text{in}}, \mathcal{B}_{\text{out}}, \mathcal{U}, \mathcal{L})$  where:*

- $\mathcal{P} = (G, \mathcal{E}, \mathcal{D}, \mathcal{R}, \mathcal{L}_{\text{MDL}})$  is a Phaneron instance (graph, event log, dictionary, rewrite schemas, MDL objective);
- $\mathcal{B}_{\text{in}}$  maps surface streams to candidate graph deltas (rewrite proposals with alignments);
- $\mathcal{B}_{\text{out}}$  maps meaning slices to surface realizations in a target language;
- $\mathcal{U}$  selects a consistent subset of proposals (constraint-driven abduction, CDA) and applies DPO rewrites;
- $\mathcal{L} = \mathcal{L}_{\text{MDL}} + \lambda_1 \mathcal{L}_{\text{pred}} + \lambda_2 \mathcal{L}_{\text{cycle}} + \lambda_3 \mathcal{L}_{\text{min}} + \lambda_4 \mathcal{L}_{\text{align}}$  is the global objective.



Figure 1: **LMM pipeline.** Surface streams are mapped to *proposals* (Bridge-In); the Phaneron applies versioned rewrites; Bridge-Out realizes responses.



Figure 2: **Architecture layers.** Surface constructions map through bridge patterns into language-agnostic meaning patterns; realization reverses the mapping.

## 4 Pattern typology for language

We use canonical unlabeled motifs for linguistic phenomena; roles are structural positions inside relation patterns.

**Events and roles.** Clausal semantics as an event node with role positions (agent, patient, adjuncts).

**Alternations.** Active/passive/dative shift map to the same meaning motif; only bridge patterns differ.

**Quantification and scope.** Restrictor–scope encoded structurally; ambiguities are competing pattern placements.

**Negation/modality/tense.** Operator patterns attach to events; scope is explicit via structure.

**Coreference and anaphora.** Entity nodes persist; pronouns propose merges with antecedent candidates.

**Presupposition and discourse.** Triggers require prior subgraphs; accommodation introduces minimal subgraphs.

## 5 Discourse as versioned rewrites

A conversation or document is a sequence of versioned events  $e_t$ ; each utterance proposes rewrites. The state  $S_t = (G_t, \mathcal{D}_t)$  evolves via CDA selection and DPO application; replay yields an explainable discourse history.

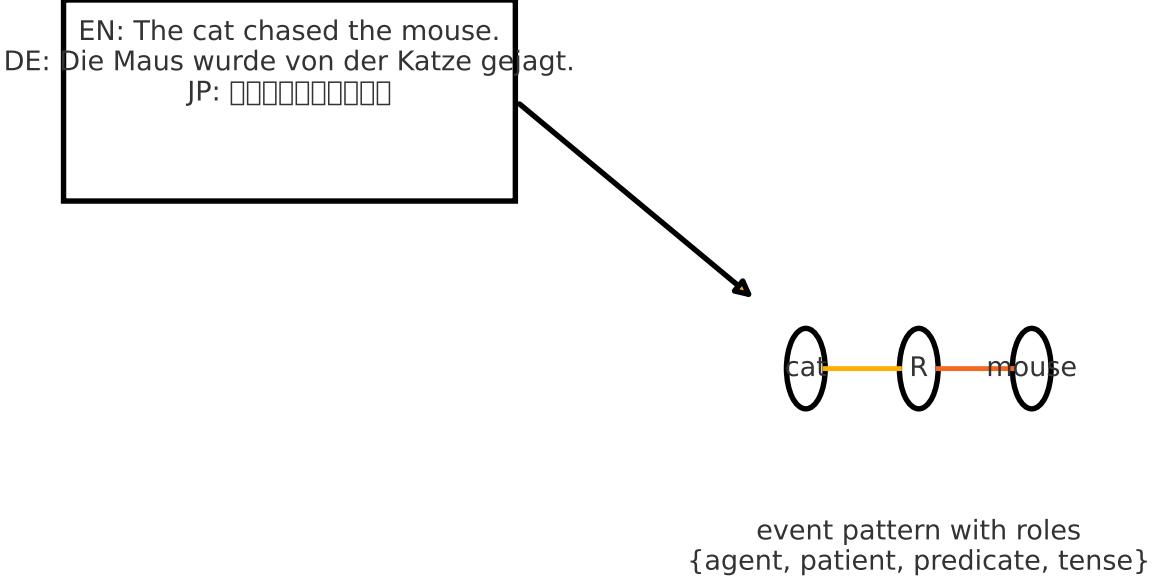


Figure 3: **Surface → meaning.** Three surface realizations (EN/DE/JP) map to a single event pattern with roles.

## 6 Bridges: proposal pathways and constraints

**Proposal pathways.** (i) Grammar-seeded parser that proposes construction-to-role alignments; (ii) LLM-assisted proposer that emits span↔role candidates; (iii) retrieval from previously accepted patterns for rapid reuse.

**CDA constraints.** Type and arity for roles; role-cardinality bounds; acyclicity where applicable; locality windows for new attachments; version-monotonic merges; minimality preference for explanations.

**Bridge-Out.** Realization chooses constructions and lexicalizations consistent with a target language and style; morphology/word order patterns fill remaining slots.

## 7 Learning objective (hybrid)

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MDL}}(G; \mathcal{D}, B) + \lambda_1 \mathcal{L}_{\text{pred}} + \lambda_2 \mathcal{L}_{\text{cycle}} + \lambda_3 \mathcal{L}_{\text{min}} + \lambda_4 \mathcal{L}_{\text{align}}. \quad (1)$$

## 8 Evaluation protocols

**Paraphrase invariance.** Multiple surface forms must map to isomorphic meaning patches (graph isomorphism/edit distance).

**Cross-lingual alignment.** Translations align to identical meaning graphs; test zero-shot transfer.

**Explainable QA.** Answers with Minimal Explanation Subgraphs (MES); evaluate correctness and sparsity.

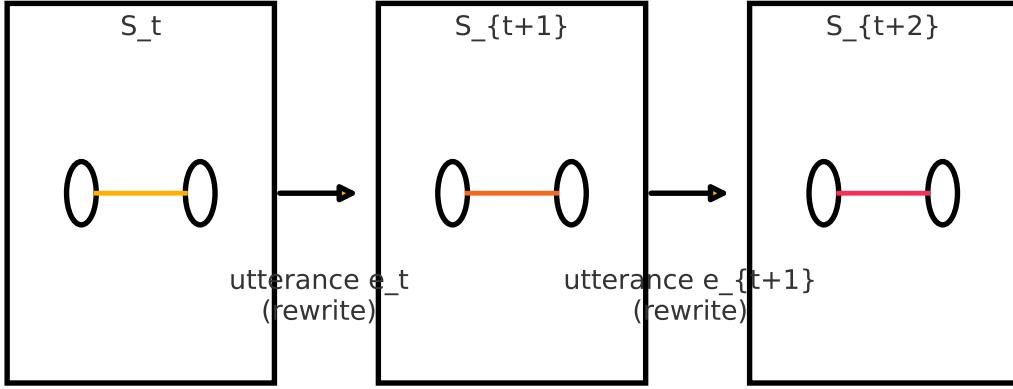


Figure 4: **Discourse timeline.** Utterances update the meaning graph by applying selected rewrites; anaphora appears as merges across versions.



Figure 5: **Inference loop.** Propose → score → select → apply.

**Discourse tracking.** Coreference and presupposition accommodation across long contexts.

Task	Metric(s)	Structural check
Paraphrase invariance	Graph isomorphism rate; edit distance	Canonicalization match; MES stability
Cross-lingual alignment	Zero-shot accuracy; alignment F1	Same meaning hash; role consistency
Explainable QA	EM / F1 + explanation F1	MES size/precision vs gold
Coreference	MUC/B <sup>3</sup> /CEAF or link F1	Merge decisions; entity continuity
Discourse relations	Accuracy / F1	Correct rhetorical link patterns
Efficiency	Latency; updates/s; growth	$ \mathcal{D} $ growth; replay time

**Evaluation plan (summary).**

## 9 Micro-bench: controlled paraphrase set

**Dataset.** 100 sentence triplets (active, passive, topicalized) per predicate, across 20 predicates; 10 languages (EN/DE/ES/FR/IT/PL/NL/SV/JA/KO) with human-checked translations.

**Split.** Train on half the predicates and 7 languages; test zero-shot on the rest and 3 held-out languages.

**Targets.** (i) Paraphrase invariance  $\uparrow$  (canonical hash equality), (ii) cross-lingual isomorphism rate, (iii) MES stability under paraphrase.

**Ablations.** No-cycle loss; no-minimality loss; fixed pattern cap vs dynamic.

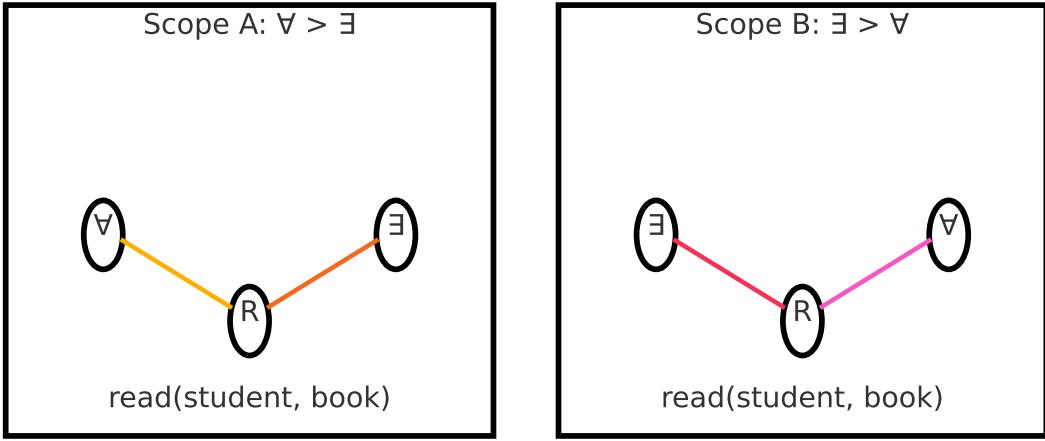


Figure 6: **Scope ambiguity.** Two competing patterns for “Every student read a book”:  $\forall > \exists$  vs  $\exists > \forall$ .

## 10 Case studies

**Quantifier scope ambiguity.** Competing structures resolved by context.

**Minimal explanation subgraph (MES).** “Alice dropped the plate. It shattered.”

**Pattern dictionary snapshot.**

## 11 A hybrid LMM+LLM system

LLMs can serve as proposal generators and realizers; the Phaneron enforces structural constraints and provides long-term semantic memory.

## 12 Limitations, risks, and security

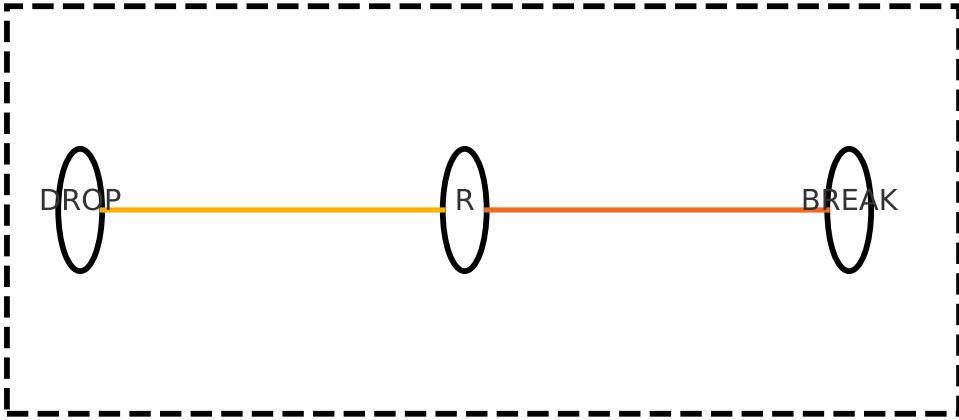
Pattern explosion, search cost, and bridge brittleness are practical risks. Explicit meaning graphs may encode sensitive facts; we recommend:

- **PII redaction:** detect and mask subgraphs linked to personal identifiers.
- **Access control:** role-based visibility of subgraph types (entities, relations).
- **Audit trails:** versioned provenance for all rewrites, with replay.
- **Safety rules:** refusal policies as subgraph pattern detectors (harmful intent).

Some pragmatic phenomena (irony, sarcasm) may require richer patterns and grounding.

## 13 Conclusion

LMMs make meaning first-class: explicit, stable under paraphrase, language-agnostic, and replayable. Phaneron provides the substrate; MDL+prediction provides the learning principle; versioned rewrites provide the discourse engine. LMMs complement LLMs: structure for memory/reasoning; fluency for realization.



Minimal Explanation Subgraph (MES): {DROP, R<sub>cause</sub>, BREAK}

Figure 7: **Explanation view.** MES = {DROP, R<sub>cause</sub>, BREAK}.

## A Appendix A: Concrete MDL code

We use a two-part code for  $(\mathcal{D}, B)$  and residual structure  $R$ . Let  $\mathcal{D} = \{P_i\}_{i=1}^M$  be unlabeled patterns;  $U_i$  their usage sets (injective embeddings).

$$L(\mathcal{D}) = \sum_{i=1}^M \left( L_{\mathbb{N}}(|V(P_i)|) + L_{\mathbb{N}}(|E(P_i)|) + L_{\text{iso}}(P_i) \right), \quad (2)$$

$$L(U \mid \mathcal{D}, B) = \sum_{i=1}^M \left( L_{\mathbb{N}}(|U_i|) + \sum_{u \in U_i} L_{\text{place}}(u \mid P_i, B) \right), \quad (3)$$

$$L(R \mid \mathcal{D}, U, B) = L_{\text{edges}}(\text{residual edges} \mid B). \quad (4)$$

Here  $L_{\mathbb{N}}$  is a universal integer code;  $L_{\text{iso}}$  encodes an unlabeled pattern up to isomorphism (canonical adjacency);  $L_{\text{place}}$  codes a placement via block-structured anchors  $B$ ; and  $L_{\text{edges}}$  uses a Bernoulli or degree-corrected block model for uncovered edges. A rewrite is admissible if total  $\Delta L$  plus predictive losses is  $< 0$ .

## B Appendix B: CDA selection (pseudocode)

```

Input: state S_t=(G_t,D_t), proposals P={ G_k }, objective L( )
1: score each G_k with local L ; discard non-improving
2: sort by L ; greedily build non-overlapping set W
3: apply DPO rewrites in W to obtain S_{t+1}; append events
4: consolidate: merge near-duplicate patterns; re-score dictionary

```

## C Appendix C: Canonicalization and tests

We canonicalize small meaning patches for fast paraphrase alignment using a hash of canonical adjacency; ties are broken by lexicographically minimal rooted traversal. Paraphrase invariance

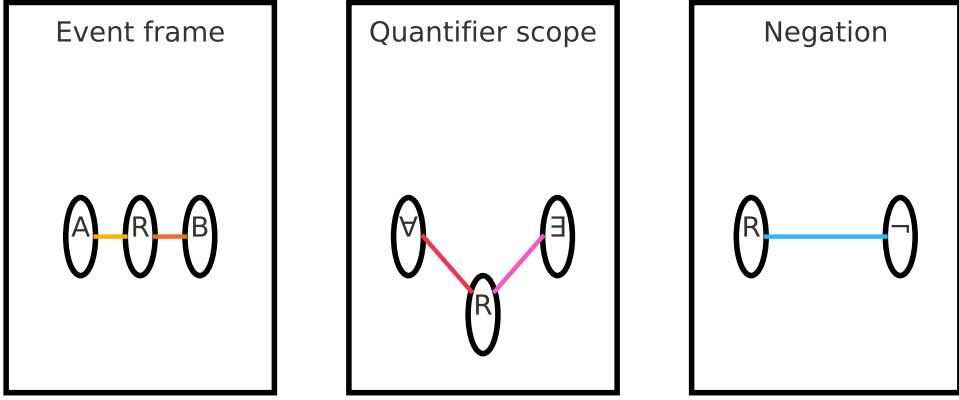


Figure 8: **Pattern dictionary motifs.** Event frame, quantifier scope, and negation attachment.



Figure 9: **Hybrid stack.** LLM front/back ends surround the explicit meaning core.

= identical canonical hashes (with audit trail).

## D References

### References

- [1] J. Rissanen. “Modeling by shortest data description.” 1978.
- [2] J. Rissanen. *Stochastic Complexity in Statistical Inquiry*. 1989.
- [3] B. Weisfeiler and A. Leman. “The reduction of a graph to canonical form and the algebra which appears therein.” 1968.
- [4] H. Ehrig et al. *Fundamentals of Algebraic Graph Transformation*. 2006.
- [5] C. Fillmore et al. “FrameNet and frame semantics.” 2003.
- [6] L. Banarescu et al. “Abstract Meaning Representation for sembanking.” 2013.
- [7] N. Abzianidze et al. “UCCA: Universal Conceptual Cognitive Annotation.” 2017.
- [8] H. Kamp and U. Reyle. *From Discourse to Logic*. 1993.
- [9] N. Asher and A. Lascarides. *Logics of Conversation*. 2003.
- [10] A. Goldberg. *Constructions: A Construction Grammar Approach*. 1995.
- [11] P. Lewis et al. “Retrieval-Augmented Generation for Knowledge-Intensive NLP.” 2020.

[12] R. Besold et al. “Neural-Symbolic Integration.” 2017.