

Constraint-Driven Abduction (CDA) – Technical Spec

Status: MVP – implementation-ready

Audience: Internal (Metis Core + Daedalus partner profile)

Container: Conversation (Canvas-first)

Transports: WebSocket ToolBus (control) · SSE (streams). **No REST anywhere.**

Contracts: ToolBus **JSON schemas** (in/out) + **SSE event names** only.

Terminology refresh: **Actions** replace Work Units (WUs). **Abilities** are action types. CDA is invoked primarily as an **Action** with ability: "cda". There is **no public cda.* ToolBus namespace** for partners (internal-only debug exists behind flags).

0) Purpose & Scope

CDA is a bounded, auditable reasoning engine for forming the **best explanatory hypotheses** under explicit **constraints** given observations and goals. It orchestrates iterative **hypothesis generation → constraint propagation → verification → plan refinement**. CDA prioritizes **safety, determinism, and boundary-only disclosure** (no raw chain-of-thought to clients).

Use cases (examples): root-cause analysis, multi-step verification of claims, plan synthesis under conflicting evidence, policy compliance checks, red-team scenario analysis.

Non-goals (MVP): unconstrained creative brainstorming (use Compose/Generative), unrestricted web crawling (use Research with approvals), public REST.

1) Mental Model

- **Inputs:** observations O , constraints C (hard/soft), goals G (prove/disprove, best explanation, feasible plan), background knowledge via **Ken**.
- **Outputs:** a **bounded hypothesis set** H^* with scores/confidence, **boundary summaries**, **verifier justifications**, and **next-best tests**.

- **Search:** seeded, deterministic exploration of hypothesis space with pruning by constraints and verification feedback.
 - **Disclosure:** only **boundary summaries** and artifact IDs are streamed to clients; raw traces stay internal behind Reflection (gated).
-

2) Interfaces

2.1 Action Inputs

```
CDAInputs {
    problem: string,                                // natural-language statement of the
    abductive task
    observations: SpanRef[] | ArtifactRef[],
    goals: { kind: "best_explanation|prove|disprove|plan", success_criteria?: string[] },
    constraints: {
        hard?: string[],                            // DSL clauses (see §3)
        soft?: {clause: string, weight?: number}[]
    },
    policies?: {pii_posture?: "on|off|consent", egress?: "local|partner|external"},
    budgets?: {tokens?: number, wall?: number},
    risk_class?: "low|medium|high|hard",      // hard => escalate immediately
    seed?: {purpose?: string}
}
```

2.2 Action Outputs

```
CDAOOutputs {
    summary_artifact_id: string,                  // boundary summary document
    plan_bundle_id?: string,                      // cda_plan (bundle) with machine-
    readable plan graph/constraints
    verdicts: Array<{hypothesis_id:string, status:"supported|refuted|inconclusive", confidence:number}>,
    next_tests?: Array<{question:string, suggested_ability?:"research|sandbox|default", expected_discriminative_power?:number}>,
}
```

```
    evidence_bundle_id?: string          // citations/SpanRefs for verifier
decisions
}
```

2.3 DisplayPlan (client hints)

```
DisplayPlan {
  primary: {kind:"summary", id_ref:"summary_artifact_id"},  

  chips: ["receipts", "context_pack"],  

  drawers: {integrity:true}  

}
```

3) Constraint DSL (MVP)

A compact, auditable DSL for **hard** and **soft** constraints. Parsed and validated server-side; logged in receipts.

Forms (allowed):

- **Boolean predicates:** holds(pred(args)), not(pred(args)).
- **Relational over Ken:** exists path(entity:A ->related_to-> entity:B within 2 hops).
- **Temporal ordering:** before(event:X, event:Y), within(event:X, 14d of event:Y).
- **Mutual exclusivity:** xor(hyp:A, hyp:B).
- **Numeric bounds:** value(metric:Cost) <= 5000.
- **Selection policy:** use(provider:ModelX) only_if risk <= medium (interpreted by Router caps).
- **Soft clause:** prefer(pred(args)) weight 0.4.

Safety: only the **safe query set** to Ken is permitted (search|neighborhood|path|related_by); any external egress triggered by a clause requires approvals-as-input. PII posture gates predicates over personal data.

4) Algorithm (Deterministic, Budget-Aware)

- 1) **Initialize** frontier from observations 0 using seeded generators (deterministic enumeration → beam of candidates).
- 2) **Propagate constraints**: apply hard constraints; compute penalties for soft constraints.
- 3) **Verify**: run verifiers (LLM fact-checkers, symbolic checkers, typed functions/tests) & **Ken** queries; attach **citations (SpanRefs)**; emit `action_checkpoint(kind:"verifier_update")`.
- 4) **Refine plan**: expand or prune hypotheses; emit `action_checkpoint(kind:"plan_delta")` with boundary-safe rationale.
- 5) **Context updates**: when adopting new evidence, emit `action_checkpoint(kind:"context_delta")` and update the active **context pack**.
- 6) **Converge**: stop when success criteria met, budget exhausted, or frontier stable; produce **boundary summary + next-best tests**.

Search control: beam width, depth caps, and reranker settings are deterministically derived from **hmac-v1** seeds and budgets.

Escalation: may spawn sub-Actions (`research`, `sandbox`, or `default`) via edges `dependsOn` | `waitFor|bind|notify`; outputs are bound as locked dependencies.

5) Verification Suite (MVP)

- **LLM verifiers**: cross-model agreement, self-consistency checks, claim-level citation enforcement.
 - **Symbolic verifiers**: constraint satisfiability, unit checks, temporal logic evaluation.
 - **Function evaluators**: deterministic code tests in **Sandbox** (egress denied unless approved).
 - **Ken evidencing**: `search|neighborhood|path|related_by` with receipts; Weave supplies span pointers.
 - **Scoring**: $\text{score} = \text{salience} + \text{posterior_support} - \text{penalties}(\text{violations}) - \text{uncertainty}$; confidence bands derived from calibration tables.
-

6) Privacy, Egress & Reflection

- **Boundary-only**: clients receive fixed-size boundary summaries; raw traces accessible only via **Reflection** (read-only, approval-gated).
- **PII posture**: `on|off|consent`; when **on**, external provider calls are redacted; concept-only

embeddings allowed with approval.

- **Egress classes:** local|partner|external; Router enforces caps and approvals.
 - **Audit/Receipts:** every verifier and Ken call logs a **receipt** with seeds, caps, and citations; WORM-logged.
-

7) Public Contracts

7.1 ToolBus (WS)

- **No public** cda.* for partners. Invoke via `actions.start{ ability:"cda", inputs:CDAInputs }`.
- **Sub-abilities:** CDA may call research.*, sandbox.*, and ken.* through the orchestrator.
- **Reflection:** reflect.* (read-only) for internal reviewers (approval-gated).

7.2 SSE Events (Streams)

- **Actions:** action_started, action_progress, action_notice, action_waiting_user, action_linked, action_checkpoint(kind:"boundary_summary|plan_delta|verifier_update|context_delta"), action_completed, action_partial, action_failed.

Cursor resume: by Last-Event-ID; ≥24h retention per topic.

8) Artifacts & Display

- **Boundary summary (primary):** short document with decisions, key evidence citations, and residual uncertainty.
- **Plan bundle (cda_plan):** **bundle** artifact containing {hypotheses.json, constraints.json, plan_graph.json, verifiers.json, receipts/} (no raw CoT).
- **Evidence bundle:** citations with SpanRefs and excerpts; missing evidence yields **receipt stubs**.
- **DisplayPlan** emphasizes the boundary summary; Plan bundle & receipts are available via chips/drawers.

9) CUP → CDA Escalation & Budgets

- **Triggers:** escalate to CDA when **any 2 of 5** uncertainty/risk signals fire (e.g., contradictions in evidence, high doubt, low confidence band, policy risk, critical decision), or **immediately** for `risk_class: "hard"`.
 - **Budget reservation:** CDA should reserve ~**20%** of wall time ($\geq 20\text{ s}$) upon escalation; if less remains, emit **partial** with explicit caveats.
 - **Grace → freeze → partials** applies on budget breach; **Incomplete** badge shown in Canvas.
-

10) Observability & SLOs

- **Metrics:** hypothesis frontier size, constraints satisfied/violated, verifier pass rate, citation coverage, Ken query p95, TTFS (time-to-first summary), partial ratio, approvals latency.
 - **Targets (MVP):** TTFS $\leq 3.0\text{ s}$ for boundary summary draft; citation coverage $\geq 95\%$ for supported claims; Ken p95 $\leq 700\text{ ms}$ @ top-50.
 - **Degradation ladder:** shrink frontier/beam \rightarrow simplify constraints \rightarrow reduce external calls \rightarrow return conservative boundary summary with next tests.
-

11) Acceptance Criteria (MVP)

- 1) CDA runs as an **Action** (`ability: "cda"`), streaming **action_checkpoint** events with kinds `boundary_summary|plan_delta|verifier_update|context_delta`.
- 2) **No public** `cda.*` for partners; all invocations go through `actions.start`.
- 3) **Boundary-only** client output; raw traces accessible via **Reflection** (approval-gated).
- 4) **Ken** integration uses only the safe set; every query and verifier call emits **receipts** with seeds and citations; Weave span pointers present or stubs recorded.
- 5) **Bind ⇒ lock** when consuming sub-Action outputs; **update-available** bump produces a **structure revision** on the consumer.
- 6) **CUP escalation** rules enforced; ~20% wall reserved; partials correctly labeled on breach.
- 7) **PII posture** and **egress caps** honored; approvals-as-input audited.
- 8) **ToolBus + SSE only**; cursor resume works across deploys.

-
- 9) SLOs in §10 met on staging.
-

12) QA Scenarios & Runbooks

- 1) **Ambiguous claim:** CDA receives conflicting observations; emits verifier_update → plan_delta cycles; produces boundary summary with confidence bands and next tests.
 - 2) **Bind bump:** CDA spawns Research Action; later revision of evidence triggers **update available**; accepting bump yields structure revision; receipts updated.
 - 3) **Denied egress:** constraint requires external check; approval denied; CDA returns conservative boundary summary and marks missing evidence with receipt stubs.
 - 4) **Budget exhaustion:** CDA reserves 20%; if consumed, emits terminal **partial** with explicit caveats and next tests.
 - 5) **Cursor loss:** SSE reconnect with `Last-Event-ID`; timeline consistent.
-

13) Compatibility & Migration (WU → Actions)

- Replace all mentions of **CDA WU** with **CDA Ability** running as an **Action**.
 - Legacy `wu_*` event consumers must switch to `action_*`; optional aliases may exist for one CalVer window (off by default for partners).
 - Artifacts referencing `work_unit_id` map to `action_id` and are preserved under `legacy.aliases[]` in receipts.
-

14) Open Questions (track to backlog)

- Constraint DSL evolution: add quantifiers/aggregation safely or keep pattern-based subset?
- Learned scoring vs rule-based calibration for confidence bands.
- Canonical schema for **next-best test** utility estimates across abilities.
- Partitioned caching of partial hypothesis frontiers for warm resumes.

#ideas/metis